

1. When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e.,  $n < 30$ ). It turns out, that how “far off” the  $t$ -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

- (a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6}}_{\text{error}} f_Z(t),$$

where  $f_Z(\cdot)$  and  $F_Z(\cdot)$  are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the  $p$ -value when testing  $H_0 : \mu_X = 0; H_a : \mu_X < 0$  using the zebra finch further data?

```
library(tidyverse)
library(e1071)
set.seed(5656)
zf <- read_csv("zebrafinches.csv")

n <- length(zf$further)
skew <- skewness(zf$further)
skew.sqn <- skew/sqrt(n)
sd.zf <- sd(zf$further)
t <- mean(zf$further) / (sd.zf / sqrt(n))
(error <- (skew.sqn)*((2*t^2+1)/6)*dnorm(t))

## [1] -1.226006e-13
```

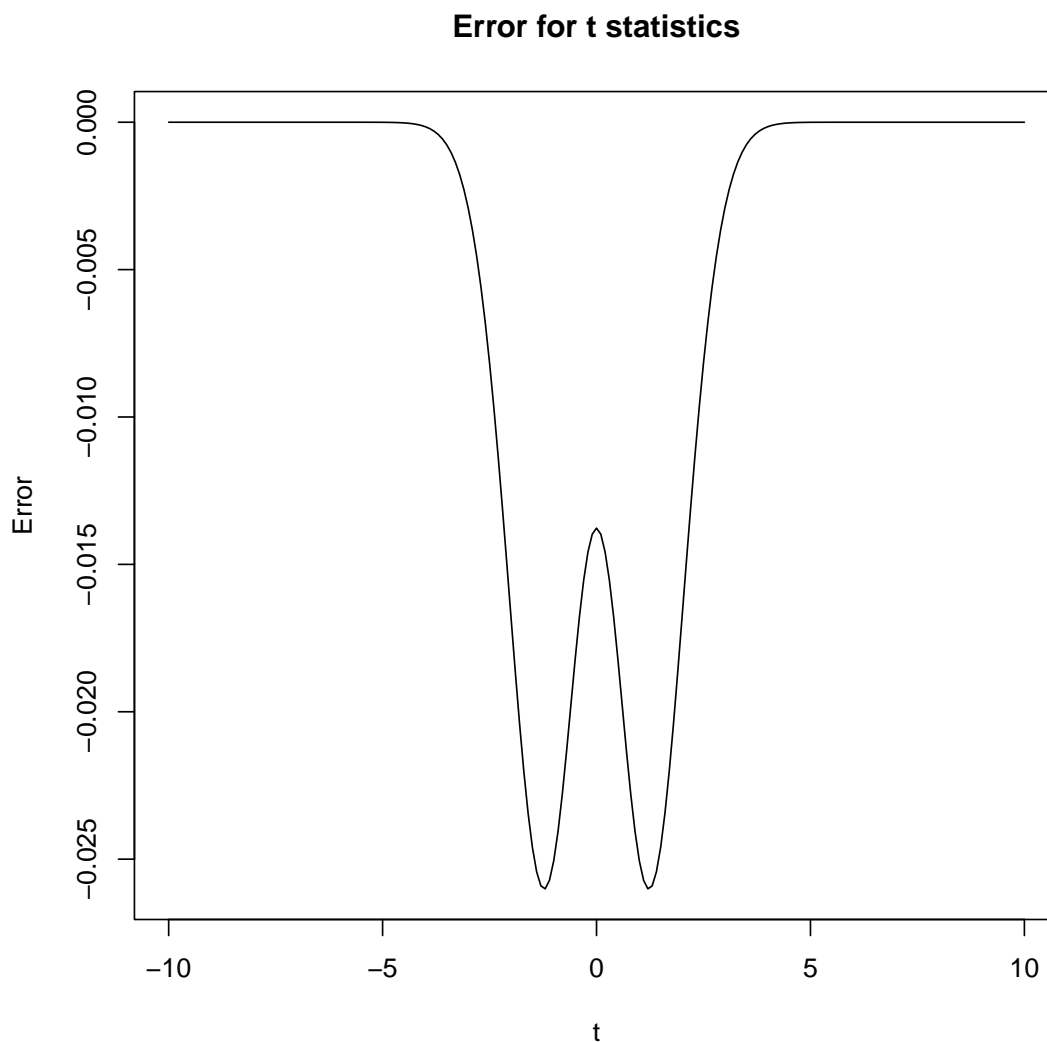
Therefore the potential error is approximately  $-1.226 \times 10^{-13}$ .

- (b) Compute the error for  $t$  statistics from -10 to 10 and plot a line that shows the error across  $t$ . Continue to use the skewness and the sample size for the zebra finch further data.

```
t.vals <- seq(-10, 10, 0.1)
errors <- tibble(error = numeric(length(t.vals)))

for(i in 1:length(t.vals)){
  errors$error[i] <- (skew.sqn)*((2*t.vals[i]^2+1)/6)*dnorm(t.vals[i])
}

plot(t.vals, errors$error, type = "l",
     main = "Error for t statistics",
     ylab = "Error", xlab = "t")
```



- (c) Suppose we wanted to have a tail probability within 10% of the desired  $\alpha = 0.05$ . Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{\text{set}}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left( \frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

```
alpha <- 0.05
t.crit <- qnorm(alpha)
(n.star <- ((skew/(6*0.10*alpha))*(2*t.crit^2+1)*dnorm(t.crit))^2)
## [1] 520.8876
```

Therefore we would need a sample size of at least 521.

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the  $T$  statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where  $\bar{x}_r$  is the mean computed on the  $r^{th}$  resample and  $s$  is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to  $t = 0$ , for each case).

```
R <- 10000
resamples <- tibble(xbars = rep(NA, R))

for(i in 1:R){
  curr.resample <- sample(zf$further,
                        size = nrow(zf),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}
r.xbar.further <- resamples$xbars

t.further <- r.xbar.further / (sd.zf / sqrt(n))
resamples.null.further <- t.further - mean(t.further)

for(i in 1:R){
  curr.resample <- sample(zf$closer,
                        size = nrow(zf),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}
r.xbar.closer <- resamples$xbars

t.closer <- r.xbar.closer / (sd.zf / sqrt(n))
resamples.null.closer <- t.closer - mean(t.closer)

for(i in 1:R){
  curr.resample <- sample(zf$diff,
                        size = nrow(zf),
                        replace = T)

  resamples$xbars[i] <- mean(curr.resample)
}
r.xbar.diff <- resamples$xbars

t.diff <- r.xbar.diff / (sd.zf / sqrt(n))
resamples.null.diff <- t.diff - mean(t.diff)
```

- (b) Compute the bootstrap  $p$ -value for each test using the shifted resamples. How do these compare to the  $t$ -test  $p$ -values?

```
t.test.further <- t.test(zf$further, mu = 0, alternative = "less")
(p.t.further <- t.test.further$p.value)

## [1] 2.587359e-08

(p.boot.further <- mean(resamples.null.further <= t.test.further$statistic))

## [1] 0

t.test.closer <- t.test(zf$closer, mu = 0, alternative = "greater")
(p.t.closer <- t.test.closer$p.value)

## [1] 8.131533e-09

(p.boot.closer <- mean(resamples.null.closer >= t.test.closer$statistic))

## [1] 0

t.test.diff <- t.test(zf$diff, mu = 0, alternative = "two.sided")
(p.t.diff <- t.test.diff$p.value)
```

```
## [1] 1.036907e-08

(p.boot.diff <- mean(resamples.null.diff <= -t.test.diff$statistic) +
mean(resamples.null.diff >= t.test.diff$statistic))

## [1] 0
```

In each case, the p-values are all very close to zero.

- (c) What is the 5<sup>th</sup> percentile of the shifted resamples under the null hypothesis? Note this value approximates  $t_{0.05, n-1}$ . Compare these values in each case.

```
(quantile(resamples.null.further, alpha))

##          5%
## -1.646715

(quantile(resamples.null.closer, alpha))

##          5%
## -1.159458

(quantile(resamples.null.diff, alpha))

##          5%
## -2.500706

(qt(alpha, n-1))

## [1] -1.710882
```

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the  $t$ -test confidence intervals?

```
t.test.further <- t.test(zf$further, mu = 0, alternative = "two.sided")
(t.test.further$conf.int)

## [1] -0.2565176 -0.1489313
## attr(,"conf.level")
## [1] 0.95

(quantile(r.xbar.further, c(0.025, 0.975)))

##          2.5%          97.5%
## -0.2558901 -0.1569692

t.test.closer <- t.test(zf$closer, mu = 0, alternative = "two.sided")
(t.test.closer$conf.int)

## [1] 0.1173875 0.1950586
## attr(,"conf.level")
## [1] 0.95

(quantile(r.xbar.closer, c(0.025, 0.975)))

##          2.5%          97.5%
## 0.1201672 0.1929522

(t.test.diff$conf.int)

## [1] 0.2719028 0.4459921
## attr(,"conf.level")
## [1] 0.95

(quantile(r.xbar.diff, c(0.025, 0.975)))

##          2.5%          97.5%
## 0.2816355 0.4414041
```

Overall, the bootstrap confidence intervals are very similar to the  $t$ -test confidence intervals.

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```

R <- 10000
rand.further <- tibble(xbars = rep(NA, R))
rand.closer <- tibble(xbars = rep(NA, R))
rand.diff <- tibble(xbars = rep(NA, R))

further.shift <- zf$further - mean(zf$further)
closer.shift <- zf$closer - mean(zf$closer)
diff.shift <- zf$diff - mean(zf$diff)

for(i in 1:R){
  curr.rand.further <- further.shift *
    sample(x = c(-1, 1),
           size = length(further.shift),
           replace = T)
  curr.rand.closer <- closer.shift *
    sample(x = c(-1, 1),
           size = length(closer.shift),
           replace = T)
  curr.rand.diff <- diff.shift *
    sample(x = c(-1, 1),
           size = length(diff.shift),
           replace = T)

  rand.further$xbars[i] <- mean(curr.rand.further)
  rand.closer$xbars[i] <- mean(curr.rand.closer)
  rand.diff$xbars[i] <- mean(curr.rand.diff)
}

rand.further <- rand.further |>
  mutate(xbars = xbars + mean(zf$further)) # shifting back
rand.closer <- rand.closer |>
  mutate(xbars = xbars + mean(zf$closer)) # shifting back
rand.diff <- rand.diff |>
  mutate(xbars = xbars + mean(zf$diff)) # shifting back

```

- (b) Compute the randomization test  $p$ -value for each test.

```

(p.rand.further <- mean(rand.further$xbars - mean(zf$further) <= mean(zf$further)))

## [1] 0

(p.rand.closer <- mean(rand.closer$xbars - mean(zf$closer) >= mean(zf$closer)))

## [1] 0

(p.rand.diff <- mean(abs(rand.diff$xbars) - mean(zf$diff) >= abs(mean(zf$diff))))

## [1] 0

```

- (c) Compute the randomization confidence interval by iterating over values of  $\mu_0$ .

**Hint:** You can “search” for the lower bound from  $Q_1$  and subtracting by 0.0001, and the upper bound using  $Q_3$  and increasing by 0.0001. You will continue until you find the first value for which the two-sided  $p$ -value is greater than or equal to 0.05.

```

R <- 1000
mu0.iterate <- 0.01
starting.point <- mean(zf$further)

mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$further - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back
}

```

```

# p-value
(delta <- abs(mean(zf$furthest) - mu.lower))
(low <- mu.lower - delta) # mirror
(high <- mu.lower + delta) # xbar
(p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high))

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$furthest - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(zf$furthest) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}
# further
c(mu.lower, mu.upper)

## [1] -0.2627244 -0.1427244

R <- 1000
mu0.iterate <- 0.01
starting.point <- mean(zf$closer)

mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$closer - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back

  # p-value
  (delta <- abs(mean(zf$closer) - mu.lower))
  (low <- mu.lower - delta) # mirror
  (high <- mu.lower + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

```

```

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$closer - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
  mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(zf$closer) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}
# closer
c(mu.lower, mu.upper)

## [1] 0.1162231 0.1962231

R <- 1000
mu0.iterate <- 0.01
starting.point <- mean(zf$diff)

mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$diff - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
  mutate(xbars = xbars + mu.lower) # shifting back

  # p-value
  (delta <- abs(mean(zf$diff) - mu.lower))
  (low <- mu.lower - delta) # mirror
  (high <- mu.lower + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.lower <- mu.lower - mu0.iterate
  }
}

```

```

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zf$diff - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
  mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(zf$diff) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}
# diff
c(mu.lower, mu.upper)

## [1] 0.2689475 0.4489475

```

Therefore the further confidence interval is  $[-0.2627244, -0.1427244]$ , the closer confidence interval is  $[0.1162231, 0.1962231]$ , and the difference confidence interval is  $[0.2689475, 0.4489475]$ .

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the  $T$  statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples ( $s_r$ ), instead of the original sample ( $s$ )

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

## References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does  $n$  have to be for  $z$  and  $t$  intervals? *The American Statistician*, 54(2):121–128.